



Analysis

Andrew Redd, PhD.

R Bootcamp 2020

Hypothesis Tests

T-test

Something Ridiculous

Test for the difference in GDP Per capita.

```
# using the restricted data from the exercise
(gdp.test <- t.test(NY.GDP.PCAP.CD~in.Africa, data=data))

##
##  Welch Two Sample t-test
##
## data:  NY.GDP.PCAP.CD by in.Africa
## t = -7.0402, df = 3.022, p-value = 0.005749
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  -58966.64 -22356.59
## sample estimates:
## mean in group Africa  mean in group Other
##           1201.708           41863.325
```

Terminology

Formulas, created with a `~`, represent relationships. They can

- be one sided `~x`
 - Lambda functions `~log(.+1)`
- or two sided `y ~ x + z`
 - specify relationships or models
- often include functions
 - `y ~ x + log(z)`
 - `y ~ x + poly(z,3)`, a polynomial fit of degree 3 on `z`
- have special syntax
 - interaction `y ~ x:z`
 - crossing `y ~ a*b` is equal to `y ~ a + b + a:b`
 - nesting `y ~ a + b %in% a` or equivalently `y~a/b`

Getting usable results from a model

The `gdp.test` object is a `htest` object, which prints nicely but what if we want to include this in our table 1?

Try these:

```
str(gdp.test) # get the underlying structure of the object.  
glimpse(gdp.test) # alternative to str that handles some objects better.  
gdp.test$p.value  
getElement(gdp.test, 'p.value')  
gdp.test[['p.value']]
```

Exercise: Proportion Test

Use `prop.test` to test for the difference in proportions of males for all countries.

5:00

Solution

```
suppressMessages(library(magrittr))
data %>%
  mutate(N.males = round(SP.POP.TOTL*SP.POP.TOTL.MA.ZS/100)) %$%
  prop.test(N.males, SP.POP.TOTL)

##
## 11-sample test for equality of proportions without
## continuity correction
##
## data: N.males out of SP.POP.TOTL
## X-squared = 135837, df = 10, p-value < 2.2e-16
## alternative hypothesis: two.sided
## sample estimates:
##   prop 1   prop 2   prop 3   prop 4   prop 5   prop 6
## 0.4781971 0.4848457 0.5014707 0.5014707 0.4999245 0.5062502
##   prop 7   prop 8   prop 9   prop 10   prop 11
## 0.4862119 0.4981671 0.4915102 0.4925855 0.4944933
```

Solution, what I did

```
suppressMessages(library(magrittr)) #< To ignore unimportant message.  
data %>%  
  mutate(N.males = round(SP.POP.TOTL*SP.POP.TOTL.MA.ZS/100)) %$%  
  prop.test(N.males, SP.POP.TOTL)
```

1. Load magrittr to get the accessor pipe %\$%.
2. mutate(N.male = ...), get counts of males, because prop.test expects counts.
3. used %\$% the Accessor pipe to make the variables directly accessible to the next command
4. used prop.test to perform the test.

Alternate Version

Passing in a matrix with 'successes' and 'failures':

```
data %>%  
  # get our 'successes' and 'failures'  
  transmute( males = round(SP.POP.TOTL*SP.POP.TOTL.MA.ZS/100)  
             , females = SP.POP.TOTL-males  
             ) %>%  
  # convert to a matrix  
  as.matrix() %>%  
  prop.test()  
  
##  
## 11-sample test for equality of proportions without  
## continuity correction  
##  
## data: .  
## X-squared = 135837, df = 10, p-value < 2.2e-16  
## alternative hypothesis: two.sided  
## sample estimates:  
##   prop 1   prop 2   prop 3   prop 4   prop 5   prop 6  
## 0.4781971 0.4848457 0.5014707 0.5014707 0.4999245 0.5062502  
##   prop 7   prop 8   prop 9   prop 10   prop 11  
## 0.4862119 0.4981671 0.4915102 0.4925855 0.4944933
```

Linear Models

Bordeaux Wine Data

```
wine <- read.csv("data/Bordeaux.csv")
```

Variables

- Wine - the name of the wine.
- Price - The price of the wine in pounds sterling (£)
- ParkerPoints - the rating out of 100 given by [Robert Parker](https://www.robertparker.com/) (<https://www.robertparker.com/>).
- CoatesPoints - the rating out of 20 given by [Clive Coates](http://www.clive-coates.com/) (<http://www.clive-coates.com/>)
- P95andAbove - a dummy variable, 1 if ParkerPoints ≥ 95
- FirstGrowth - indicator of if the wine is a [first growth](https://en.wikipedia.org/wiki/First_Growth) (https://en.wikipedia.org/wiki/First_Growth)
- CultWine - indicator of if the wine is a [cult wine](https://en.wikipedia.org/wiki/Cult_wine) (https://en.wikipedia.org/wiki/Cult_wine)
- Pomerol - indicator for if the wine is from [Pomerol](https://en.wikipedia.org/wiki/Pomerol) (<https://en.wikipedia.org/wiki/Pomerol>), France
- VintageSuperstar - indicator if the wine is a superstar

Wine Model

We will rely on the normal approximation for proportions.

```
model <- lm( Price ~ . - Wine - P95andAbove, data=wine)
```

```
model
```

```
##
```

```
## Call:
```

```
## lm(formula = Price ~ . - Wine - P95andAbove, data = wine)
```

```
##
```

```
## Coefficients:
```

```
##      (Intercept)      ParkerPoints      CoatesPoints
```

```
##      -7390.78           61.94           116.27
```

```
##      FirstGrowth      CultWine      Pomerol
```

```
##      2001.41          4583.54          739.16
```

```
## VintageSuperstar
```

```
##      1424.58
```

Not really useful.

Wine Model

We will rely on the normal approximation for proportions.

```
model <- lm( Price ~ . - Wine - P95andAbove, data=wine)
model
```

Concept

Formula Subtraction

Price ~ . - Wine - P95andAbove should be read as

*“Model Price by all variables **except** Wine and P95andAbove.”*

Summarizing Models

```
(model.summary <- summary(model))
```

```
##
## Call:
## lm(formula = Price ~ . - Wine - P95andAbove, data = wine)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3251.2  -102.5    40.5   167.2  4463.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -7390.78    3766.82  -1.962   0.0540 .
## ParkerPoints     61.94     44.84   1.381   0.1719
## CoatesPoints    116.27    119.14   0.976   0.3327
## FirstGrowth    2001.41    410.38   4.877 7.28e-06 ***
## CultWine       4583.54    470.80   9.736 2.56e-14 ***
## Pomerol        739.16    305.71   2.418   0.0184 *
## VintageSuperstar 1424.58    715.89   1.990   0.0508 .
## ---
## Signif. codes:
## 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 941.3 on 65 degrees of freedom
## Multiple R-squared:  0.7716, Adjusted R-squared:  0.7506
## F-statistic: 36.61 on 6 and 65 DF,  p-value: < 2.2e-16
```

Exercise: Try the following

Extracting Parts of the model

```
coef(model)
coef(summary(model))
deviance(model)
formula(model)
residuals(model)
```


Operations on models

```
summary(model)
plot(model)
predict(model)
vcov(model)
anova(model)
aov(model)
```

5:00

Residual plots

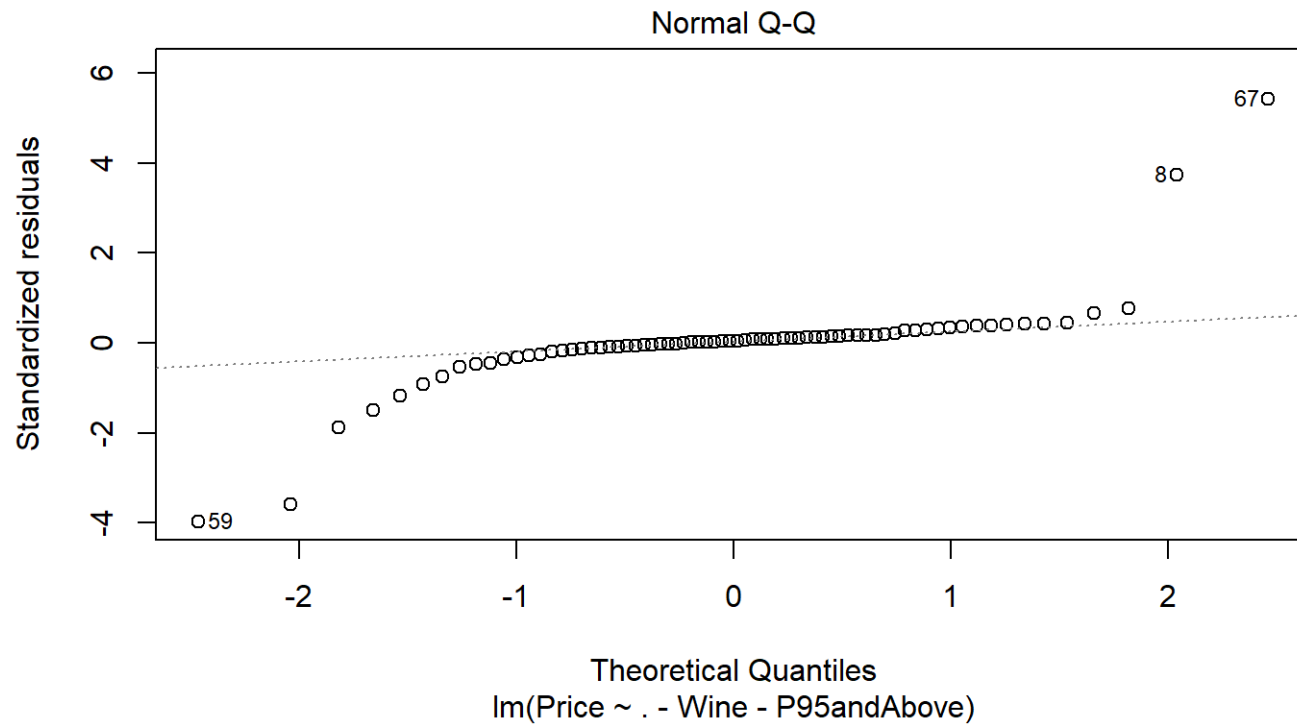
```
plot(model, 1) #< Residuals vs. Fitted | Goodness of fit
```



Residual plots

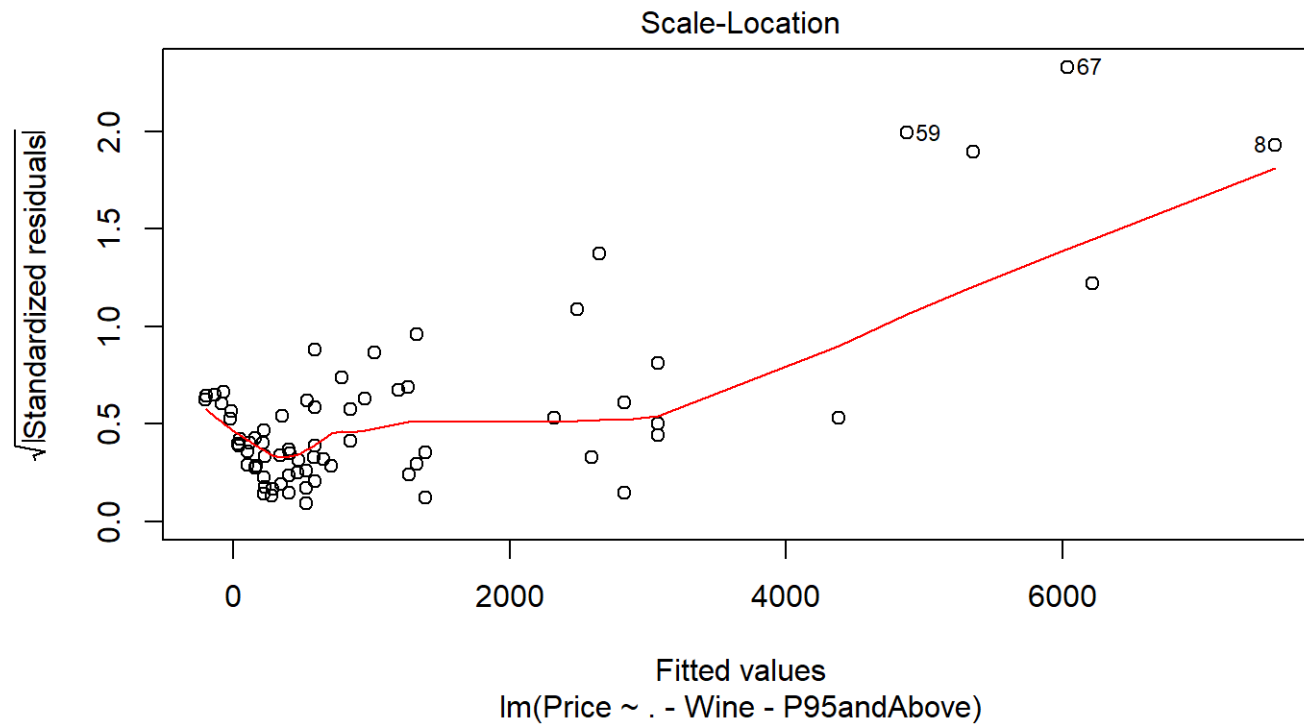
`plot(model, 2) #< Normal Q-Q`

/ Appropriate error model



Residual plots

`plot(model, 3) #< Scale-Location` */ Homoscedasticity*



Residual plots

```
plot(model, 5) #< Residual vs. Leverage | Influential points
```

